

Name: Tejas C K

SRN: PES2UG23CS647

Section: J

Email ID: tejasck05@gmail.com / pes2ug23cs647@pesu.pes.edu

Mobile Number: 9448081972

Github Link: <https://github.com/Tejas-CK/Gen-AI-Hands-on---PES2UG23CS647.git>

Assignment 1: Screenshots

```
[13] #PES2UG23CS647
!pip install transformers torch sentencepiece -q

[14] ✓ 0s
from transformers import pipeline
import warnings
warnings.filterwarnings("ignore")

[15] #PES2UG23CS647
#Text Generation pipelines
gen_bert = pipeline("text-generation", model="bert-base-uncased")
gen_roberta = pipeline("text-generation", model="roberta-base")
gen_bart = pipeline("text-generation", model="facebook/bart-base")

#Fill-Mask pipelines
mask_bert = pipeline("fill-mask", model="bert-base-uncased")
mask_roberta = pipeline("fill-mask", model="roberta-base")
mask_bart = pipeline("fill-mask", model="facebook/bart-base")

#Question Answering pipelines
qa_bert = pipeline("question-answering", model="bert-base-uncased")
qa_roberta = pipeline("question-answering", model="roberta-base")
qa_bart = pipeline("question-answering", model="facebook/bart-base")

...
-- If you want to use `BertLMHeadModel` as a standalone, add `is_decoder=True`.
Device set to use cpu
-- If you want to use `RobertaLMHeadModel` as a standalone, add `is_decoder=True`.
Device set to use cpu
Some weights of BartForCausalLM were not initialized from the model checkpoint at facebook/bart-base and are newly initialized: ['lm_head.weight', 'model.decoder.embed_tokens.weight']
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.
Device set to use cpu
Some weights of the model checkpoint at bert-base-uncased were not used when initializing BertForMaskedLM: ['bert.pooler.dense.bias', 'bert.pooler.dense.weight', 'cls.seq_relationship.bias', 'cls.seq_relationship.weight']
- This IS expected if you are initializing BertForMaskedLM from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPretraining model).
- This IS NOT expected if you are initializing BertForMaskedLM from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).
Device set to use cpu
Device set to use cpu
Device set to use cpu
Some weights of BertForQuestionAnswering were not initialized from the model checkpoint at bert-base-uncased and are newly initialized: ['qa_outputs.bias', 'qa_outputs.weight']
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.
Device set to use cpu
Some weights of RobertaForQuestionAnswering were not initialized from the model checkpoint at roberta-base and are newly initialized: ['qa_outputs.bias', 'qa_outputs.weight']
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.
Device set to use cpu
Some weights of BartForQuestionAnswering were not initialized from the model checkpoint at facebook/bart-base and are newly initialized: ['qa_outputs.bias', 'qa_outputs.weight']
You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.
Device set to use cpu
```

▼ TEXT GENERATION

▼ MASKED LANGUAGE MODELING

```
(7) #IPES2UG23CS647
mask_bert("The goal of Generative AI is to [MASK] new content.")

[{'score': 0.5396932363510132,
 'token': 3443,
 'token_str': 'create',
 'sequence': 'the goal of generative ai is to create new content.'},
 {'score': 0.15575720369815826,
 'token': 9699,
 'token_str': 'generate',
 'sequence': 'the goal of generative ai is to generate new content.'},
 {'score': 0.05405500903725624,
 'token': 3965,
 'token_str': 'produce',
 'sequence': 'the goal of generative ai is to produce new content.'},
 {'score': 0.04451530799269676,
 'token': 4503,
 'token_str': 'develop',
 'sequence': 'the goal of generative ai is to develop new content.'},
 {'score': 0.01757744885981083,
 'token': 5587,
 'token_str': 'add',
 'sequence': 'the goal of generative ai is to add new content.}]
```



```
(8) ✓ 0s
mask_roberta("The goal of Generative AI is to <mask> new content.")

[{'score': 0.3711312413215637,
 'token': 5368,
 'token_str': ' generate',
 'sequence': 'The goal of Generative AI is to generate new content.'},
 {'score': 0.3677145540714264,
 'token': 1045,
 'token_str': ' create',
 'sequence': 'The goal of Generative AI is to create new content.'},
 {'score': 0.08351420611143112,
 'token': 8286,
 'token_str': ' discover',
 'sequence': 'The goal of Generative AI is to discover new content.'},
 {'score': 0.021335121244192123,
 'token': 465,
 'token_str': ' find',
 'sequence': 'The goal of Generative AI is to find new content.'},
 {'score': 0.016521666198968887,
 'token': 694,
 'token_str': ' provide',
 'sequence': 'The goal of Generative AI is to provide new content.}]
```



```
(9) ✓ 0s
mask_bart("The goal of Generative AI is to <mask> new content.")

[{'score': 0.07461541891098022,
 'token': 1045,
 'token_str': ' create',
 'sequence': 'The goal of Generative AI is to create new content.'},
 {'score': 0.06571870297193527,
 'token': 244,
 'token_str': ' help',
 'sequence': 'The goal of Generative AI is to help new content.'},
 {"score": 0.0608800109667778015,
 'token': 694,
 'token_str': ' provide',
 'sequence': 'The goal of Generative AI is to provide new content.'},
 {"score": 0.03593561053276062,
 'token': 3155,
 'token_str': ' enable',
 'sequence': 'The goal of Generative AI is to enable new content.'},
 {"score": 0.03319477662444115,
 'token': 1477,
 'token_str': ' improve',
 'sequence': 'The goal of Generative AI is to improve new content.}]
```

QUESTION ANSWERING

```
[10] #PES2UG23CS647
qa_bert({
    "question": "What are the risks?",
    "context": "Generative AI poses significant risks such as hallucinations, bias, and deepfakes."
})

... {'score': 0.017983690835535526,
'start': 46,
'end': 82,
'answer': 'hallucinations, bias, and deepfakes.'}

[11] ✓ 0s
qa_roberta({
    "question": "What are the risks?",
    "context": "Generative AI poses significant risks such as hallucinations, bias, and deepfakes."
})

... {'score': 0.004589917603880167,
'start': 43,
'end': 82,
'answer': 'as hallucinations, bias, and deepfakes.'}

[12] ✓ 0s
qa_bart({
    "question": "What are the risks?",
    "context": "Generative AI poses significant risks such as hallucinations, bias, and deepfakes."
})

... {'score': 0.0638936311006546, 'start': 72, 'end': 81, 'answer': 'deepfakes'}
```

Task	Model	Classification (Success/Failure)	Observation (What actually happened?)	Why did this happen? (Architectural Reason)
Generation	BERT	Failure	Model failed or produced meaningless output.	BERT is an encoder-only model and cannot autoregressively generate text.
Generation	RoBERTa	Failure	Similar failure; unable to continue the sentence.	RoBERTa is also encoder-only and not trained for text generation.
Generation	BART	Success	Generated a coherent continuation of the prompt.	BART uses an encoder-decoder architecture trained for sequence generation.
Fill-Mask	BERT	Success	Correctly predicted words like "create" or "generate".	BERT is trained using Masked Language Modeling (MLM).
Fill-Mask	RoBERTa	Success	Accurate and confident predictions for the masked word.	RoBERTa improves MLM training with more data and no NSP objective.
Fill-Mask	BART	Partial Success	Filled the mask but with less precision.	BART supports masking but is optimized for seq2seq tasks rather than MLM.
QA	BERT	Partial Failure	Returned incomplete or awkward answer.	Base BERT is not fine-tuned on QA datasets like SQuAD.
QA	RoBERTa	Partial Failure	Slightly better than BERT but still inconsistent.	Encoder-only model without QA fine-tuning.
QA	BART	Success	Produced the most coherent answer.	Encoder-decoder design helps in span generation tasks.

Assignment 2: Screenshots

Character Bio Generator - Zarathos

```
[1] ✓ 7s
#PES2UG23CS647
!pip install transformers torch -q

[2] ✓ 41s
from transformers import pipeline
import warnings
warnings.filterwarnings("ignore")

WARNING:torchao.kernel.intmm:Warning: Detected no triton, on systems without Triton certain kernels will not work
```

Initialize Text Generation Pipeline

```
[3] ✓ 18s
generator = pipeline(
    "text-generation",
    model="gpt2"
)

config.json: 100% [██████████] 665/665 [00:00<00:00, 39.0kB/s]
model.safetensors: 100% [██████████] 548M/548M [00:11<00:00, 41.9MB/s]
generation_config.json: 100% [██████████] 124/124 [00:00<00:00, 11.0kB/s]
tokenizer_config.json: 100% [██████████] 26.0/26.0 [00:00<00:00, 2.66kB/s]
vocab.json: 100% [██████████] 1.04M/1.04M [00:00<00:00, 16.0MB/s]
merges.txt: 100% [██████████] 456k/456k [00:00<00:00, 7.92MB/s]
tokenizer.json: 100% [██████████] 1.36M/1.36M [00:00<00:00, 14.6MB/s]

Device set to use cpu
```

```
[4] ✓ 0s #PES2UG23CS647
prompt = (
    "Zarathos is a legendary character in a dark fantasy world. "
    "Born under a cursed eclipse, Zarathos grew up to become"
)
```

Generating Backstory

```
[5] ✓ 18s
output = generator(
    prompt,
    max_length=200,
    num_return_sequences=1,
    temperature=0.9,
    top_p=0.95
)

print(output[0]["generated_text"])

Truncation was not explicitly activated but `max_length` is provided a specific value, please use `truncation=True` to explicitly truncate examples to max length. Defaulting to 'longest_first' truncation.
Setting `pad_token_id` to `eos_token_id` to 50256 for open-end generation.
Both `max_new_tokens` (-256) and `max_length`(-200) seem to have been set. `max_new_tokens` will take precedence. Please refer to the documentation for more information. (https://huggingface.co/docs/)
Zarathos is a legendary character in a dark fantasy world. Born under a cursed eclipse, Zarathos grew up to become a human god who has turned his whole world upside down, a world where magic is forbidden.

When an accident causes the city of Zarathos to collapse, his soul is put into the hands of a mysterious demon king named Arthas, who forces Zarathos and the rest of the world to flee the planet. But
For the first time since his birth, Arthas appears. He gives Zarathos a gift, but he doesn't know what it is. But Arthas uses his power to change Zarathos into a mindless Demon Lord. He transforms Zarathos into a powerful Demon King.

In order to deal with Arthas' power, Zarathos decides to become a god. He becomes the Demon King himself, and as he
```

```
[7] ✓ 0s #PES2UG23CS647
#Cleaner Formatting
story = output[0][["generated_text"]]

print(".Character Bio: Zarathos\n")
print(story)

... Character Bio: Zarathos

Zarathos is a legendary character in a dark fantasy world. Born under a cursed eclipse, Zarathos grew up to become a human god who has turned his whole world upside down, a world where magic is forbidden.

When an accident causes the city of Zarathos to collapse, his soul is put into the hands of a mysterious demon king named Arthas, who forces Zarathos and the rest of the world to flee the planet. But
For the first time since his birth, Arthas appears. He gives Zarathos a gift, but he doesn't know what it is. But Arthas uses his power to change Zarathos into a mindless Demon Lord. He transforms Zarathos into a powerful Demon King.

In order to deal with Arthas' power, Zarathos decides to become a god. He becomes the Demon King himself, and as he
```